

Machine Learning Assignment 1

Machine Learning Package:

Package Used: Scikit Learn

I have used **Scikit Learn** for implementing the given task chosen because of the following reason:

- It's easy to use even for beginners – and a great choice for simpler data analysis tasks.
- Can understand how the data is being analysed and users have great control over data to get different outcomes.
- Also, it will be good start for me to work on ML tasks with python coding involved in it.

Features [a]:

- **Consistency:** All objects share a common interface drawn from a limited set of methods, with consistent documentation.
- **Inspection:** All specified parameter values are exposed as public attributes.
- **Limited object hierarchy:** Only algorithms are represented by Python classes; datasets are represented in standard formats (NumPy arrays, Pandas DataFrames, SciPy sparse matrices) and parameter names use standard Python strings.
- **Composition:** Many machine learning tasks can be expressed as sequences of more fundamental algorithms, and Scikit-Learn makes use of this wherever possible.
- **Sensible defaults:** When models require user-specified parameters, the library defines an appropriate default value

Data Preparation:

Data preparation before feeding data into algorithm:

1. Copy hazelnut.txt data into **hazelnut.csv**. **Transpose the data** (to correctly arrange rows and columns for the headers provided)
2. **Add headers** to each column in the order as specified in the assignment1.pdf document.

Data preparation after feeding data into algorithm:

1. **Read csv** file and store it in a data frame (using pandas).
2. **Drop 'sample_id'** (since it contains only unique ID and doesn't contribute in deciding the target) and **'variety'** (since it the target) column, and store the remaining data into X (now X contains attributes)
3. Load the values of column **'variety' into 'y'** (y contains the target data)
4. Code Implementation:

```
import pandas as pd
hazelnut = pd.read_csv('./hazelnut.csv')
print(hazelnut.head())
X = hazelnut.drop(['sample_id', 'variety'], axis=1)
y = hazelnut['variety'].values
```

References: [a]: <https://jakevdp.github.io/PythonDataScienceHandbook/05.02-introducing-scikit-learn.html>

Classification Algorithms:

Algorithm 1: K Nearest Neighbors (k-NN) ^[1]

- The principle behind nearest neighbor method is to find a predefined number of training samples closest in distance to the new point and predict the label from these.
- k-NN is a type of instance-based learning, or lazy learning, where the function is only approximated locally, and all computation is deferred until classification.
- The **input consists of the k closest training examples** in the feature space.
- The **output is a class membership**.
- An object is **classified by a plurality vote of its neighbors**, with the object being assigned to the class most common among its k nearest neighbors (k is a positive integer, typically small).
- If $k = 1$, then the object is simply assigned to the class of that single nearest neighbor.
- Choosing the **optimal value for K is best done by first inspecting the data**.
- In general, a large K value is more precise as it reduces the overall noise but there is no guarantee.
- **Cross-validation is another way to retrospectively determine a good K value** by using an independent dataset to validate the K value.
- Historically, the optimal K for most datasets has been between 3-21.

Algorithm 2: Support Vector Machines (SVM) ^{[2][3]}

- The objective of the support vector machine algorithm is to **find a hyperplane** in an N-dimensional space (N - the number of features) that distinctly classifies the data points.
- An SVM model is a **representation of the examples as points in space**, mapped so that the examples of the separate categories are **divided by a clear gap that is as wide as possible**.
- A good separation is achieved by the hyperplane that has the largest distance to the nearest training-data point of any class (so-called functional margin), since in general the **larger the margin, the lower the generalization error of the classifier**.
- New examples are then mapped into that same space and predicted to belong to a category based on the side of the gap on which they fall.
- An SVM is a discriminative classifier formally defined by a separating hyperplane.
- In other words, given labelled training data (supervised learning), the algorithm outputs an optimal hyperplane which categorizes new examples.
- In **two-dimensional space this hyperplane is a line** dividing a plane in two parts where in each class lay in either side.

Machine Learning Assignment 1




image source:

<https://towardsdatascience.com/support-vector-machine-introduction-to-machine-learning-algorithms-934a444fca47>

References:

[1]: https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm

[2]: https://en.wikipedia.org/wiki/Support-vector_machine


[3]: <https://medium.com/machine-learning-101/chapter-2-svm-support-vector-machine-theory-f0812effc72>

10-fold cross-validation:

Algorithm 1: K Nearest Neighbors (k-NN)

- To perform 10-fold cross validation, the value of k is taken as 3
- For each value fold in 10-fold cross validation, and the prediction accuracy is stored in list.
- Below code shows the implementation of 10-fold cross validation and the corresponding graph.
- Implementation takes the complete data prepared in previous step.

```
from sklearn.model_selection import cross_val_score
#create a KNN model
knn_cv = KNeighborsClassifier(n_neighbors=3)
#train model with cv of 10
cv_scores = cross_val_score(knn_cv, X, y, cv=10)
print("Cross Validation scores: \n",cv_scores)
# plot the graph for each fold in knn
cv_range = (1,2,3,4,5,6,7,8,9,10)
plt.plot(cv_range,cv_scores)
plt.xlabel('Number of folds')
plt.ylabel('10-fold cross validation accuracy')
```



Machine Learning Assignment 1

10-fold Cross Validation score for k=3:

```
print("Cross Validation scores: \n",cv_scores)
```

```
Cross Validation scores:  
[0.9047619  0.76190476 0.9047619  0.85714286 0.85714286 0.9  
0.73684211 0.94736842 0.78947368 0.78947368]
```

By taking the mean of 10-fold Cross Validation score, accuracy of the model can be measured as below:

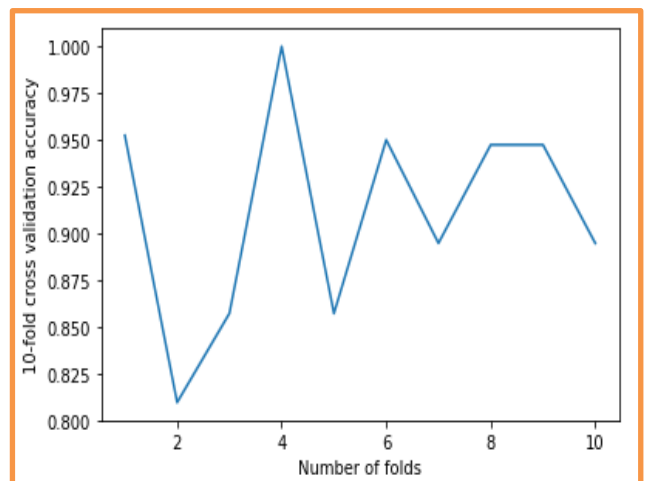
```
mean = format(np.mean(cv_scores))  
mean = float(mean)*int(100)  
print("Accuracy of the model after 10-fold cross validation: %0.2f%%mean,"%mean,"%")
```

```
Accuracy of the model after 10-fold cross validation: 84.49 %
```

Algorithm 2: Support Vector Machines (SVM)

- To perform 10-fold cross validation, **linear kernel** is used.
- A penalty parameter is introduced for the error term
- Below code shows the implementation of 10-fold cross validation and the corresponding graph.
- Implementation takes the complete data prepared in previous step.

```
import matplotlib.pyplot as plt  
from sklearn.svm import SVC  
from sklearn.model_selection import cross_val_score  
# select kernel as linear with penalty parameter C  
clf = SVC(kernel='linear', C=1)  
# train model with cv of 10  
cv_scores = cross_val_score(clf, X, y, cv=10)  
# plot the graph for each fold in knn  
cv_range = (1,2,3,4,5,6,7,8,9,10)  
plt.plot(cv_range,cv_scores)  
plt.xlabel('Number of folds')  
plt.ylabel('10-fold cross validation accuracy')
```



Machine Learning Assignment 1

10-fold Cross Validation score:

```
print("Cross Validation scores: \n",cv_scores)
```

```
Cross Validation scores:  
[0.95238095 0.80952381 0.85714286 1.          0.85714286 0.95  
 0.89473684 0.94736842 0.94736842 0.89473684]
```

By taking the mean of 10-fold Cross Validation score, accuracy of the model can be measured as below:

```
mean = cv_scores.mean()  
mean = float(mean) * int(100)  
print("Accuracy: %0.2f" %mean, "%")
```

```
Accuracy: 91.10 %
```

Conclusion:

The results and observations show that SVMs are a more reliable classifier than k-NN.

Algorithm	Training data	Accuracy
k-NN	201	84.45
SVM	201	91.10

- *KNN is less computationally intensive than SVM.*
- *Since, k-NN is easy to implement, the classification of Multi-class data should be done with k-NN.*
- *If data is homogenous to look at, one might be able to classify better by putting in a kernel into the SVM.*
- *For most practical problems, KNN is a bad choice because it scales badly, if there are a million labelled examples, it would take a long time (linear to the number of examples) to find K nearest neighbours*